

Storing Big Data

RDBMS

- Lots of work/research for the past 40 years
- Mostly centralized model
- Different cost model than in the past
- Different paradigm than most programming languages
- Provide a lot of guarantees

40 years ago : Hardware expensive
ingeni : cheap

↳ SQL vs rest

ACID

- Andreas Reuter & Theo Härder in 1983.
- Atomicity → A transaction will either succeed or fail completely
- Consistency → A transaction will move the DB from a valid state to another state
- Isolation → Many transactions can run in parallel and give the same result as a sequential execution.
- Durability → A crash should not lead to data loss

What now ?

- More data (like really more)
 - Not all well structured/organized
- Cheap hardware and not so cheap engineers
 - Many machines, 1 engineer
 - The network is everywhere
 - Machines or network **will** fail
- Is ACID possible in a distributed environment ?
 - Not really (Brewer)

BASE

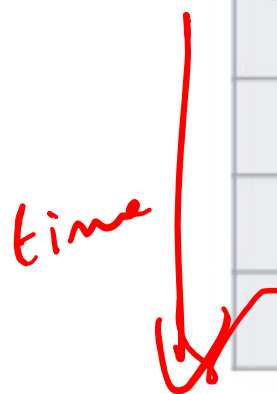
- ACID is very hard in a distributed environment
- Move to less strict guarantees
- **Basically Available** → always get an answer, even outdated
- **Soft-State** → without new requests, the state of the DS can change
- **Eventual Consistency**
↳ ?

What is consistency

- A contract between a database and a programmer
 - Follow some rules and your data will be consistent
- Many different models
 - Strict, sequential, causal, eventual...
 - Ordered from strong to weak

Strict consistency


- A write to a variable is instantaneously seen by all processors



Sequence	Strict model		Non-strict model	
	P1	P2	P1	P2
1	W(x)1		W(x)1	
2		R(x)1		R(x)0
3				R(x)1

Atomic Consistency

- Operations are executed in the same order on all machines
 - Uses a global clock
 - Same order as they were emitted
- Always deterministic




Sequential Consistency

- Weaker than strict consistency
- All write operations by multiple processors have to be seen in the same order
 - No specific order initially
 - Not necessarily consistent between various executions
- Sequential consistency + time \Rightarrow atomic consistency (e.g Google Spanner)

Eventual Consistency

- A form of weak consistency
 - Given enough time without update, all read access to a variable will return the latest value.



NoSQL databases


Principles

- Not Only SQL
- All follow the BASE principles
- Provides various properties under CAP
- Designed to scale horizontally
- Replication
 - Data is copied on multiples machines
- Various designs

Key-Value

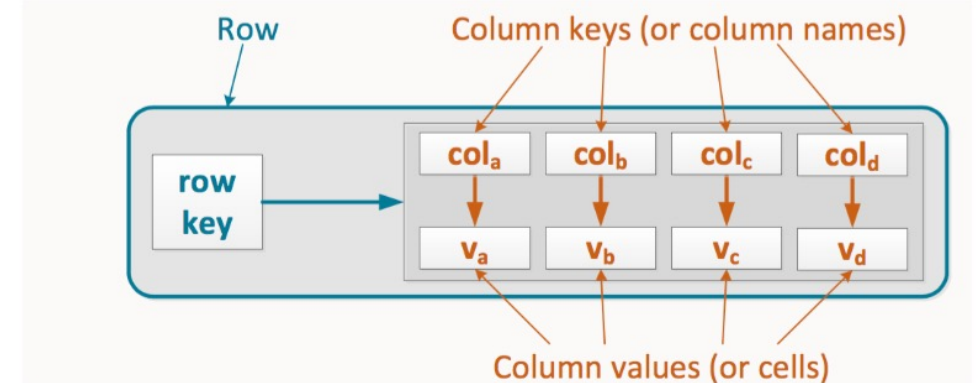
- Data are stored as unique key-value pairs
- Very simple API
 - Get, put, delete
 - Range queries often not supported
- Usually relies on consistent hashing
 - Spread keys among multiples machines
 - Copy pairs for redundancy
- Examples : DynamoDB, Redis, Riak

keys are in an interval
Divide the key space into
intervals. Assign an interval
to a machine



Wide Column

- Use row/columns to store data
 - Like RDBMS except columns have usually no fixed type
 - Number of columns can vary from row to row
- Can be seen as a 2D key-value store
(key, (key₁, v₁), (key₂, v₂))
- Examples : Apache Hbase, Cassandra



Document

- Data are stored as documents (XML, JSON...)
 - Rich data structures
 - Support versioning
- An API allows complex queries

```
db.users.insertOne(  ← collection
{
  name: "sue",        ← field: value
  age: 26,             ← field: value
  status: "pending"   ← field: value
}                    } document
)
```

- Examples : CouchDB, MongoDB


```
db.users.find(
  { age: { $gt: 18 } },
  { name: 1, address: 1 }
).limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

Graph Oriented


- Consider data as graphs
 - Introduce relations more complex than key-value

Key-value



- Examples : Neo4J, RedisGraph

Key-value as graph



Distributed Storage


Storing without DB

- Not all files require a DB
- How to store arbitrary data?
 - Hard drive
- A single hard drive is limited
 - So use many!
- Single node or multiple nodes?

Single node




- Add as many HD as possible
 - 24, 36 depending on the hardware
- Users don't like managing independent HD
 - Aggregate them to give the illusion of a single drive
 - At the hardware or OS layer
- Files are stored as blocks (4MB on recent drives)
 - Spread blocks among disks
 - RAID 0
- Pros
 - User friendly, very fast
- Cons
 - Not reliable



Single node


$$A = [A_1, A_2, A_3, A_4]$$
$$A_p : f(A_1, A_2, A_3, A_4)$$

- How to protect against data loss?
 - Add redundancy information to rebuild missing blocks
 - Use Hamming Codes
- RAID 5 or 6
 - Group of blocks + 1 or 2 parity block
 - All blocks spread among all disks
- Pros:
 - User friendly, fast
 - Reliable
- Cons
 - Loss of disk space




Multiple nodes

- Single node
 - Can be expensive
 - Is a Single Point of Failure (network outage, fire...)
- Use multiple machines
 - Each machine can use RAID
 - Add a layer on top of it
- Examples : Ceph, GlusterFS, HDFS...



Storing files on multiple nodes

- How to find blocks ?
 - Cannot query all nodes
- Metadata *data about data*
 - Added data by the OS or the filesystem
 - Date (creation, modification...), ownership..., location of blocks
- Usually “well known” nodes act as metadata servers
 - Extremely important nodes
 - Redundancy is mandatory for reliability

Distributed Storage

Case study : The Hadoop File System

Principles

- Distributed filesystem over multiple nodes
- Provides a separate and global filesystem
 - Unix like paths
 - E.g. /home on HDFS is not /home of the machine
- Not part of the OS, added software
 - Written in Java, works on any machine with JVM support
 - Shipped with script for users and administrator

Commands

- All commands are done using bin/hdfs or bin/hadoop
 - hdfs <command> [options]
 - hadoop <command> [options]
- 3 types : client, admin and daemon
- Client commands : use the filesystem
 - *hdfs dfs <args>* or *hadoop fs <args>* : run the command args on the HDFS filesystem
 - Examples : -mkdir, -put, -copyfromlocal, -get, -copytolocal, -rmr


Commands

- Admin commands : manage the filesystem
 - `hdfs fsck` : check and repair the filesystem
- Daemon commands : manage the infrastructure
 - `hdfs balancer` : run the balancy utility
 - `hdfs datanode` : start a new datanode
 - `hdfs namenode` : start a new namenode

↳ Nodes

Architecture

- Data and metadata are separated
 - Data are stored in DataNodes
 - Metadata are stored in NameNodes
- Data are divided into blocks
 - Default value 64MB or 128MB
- Blocks are stored on different machines

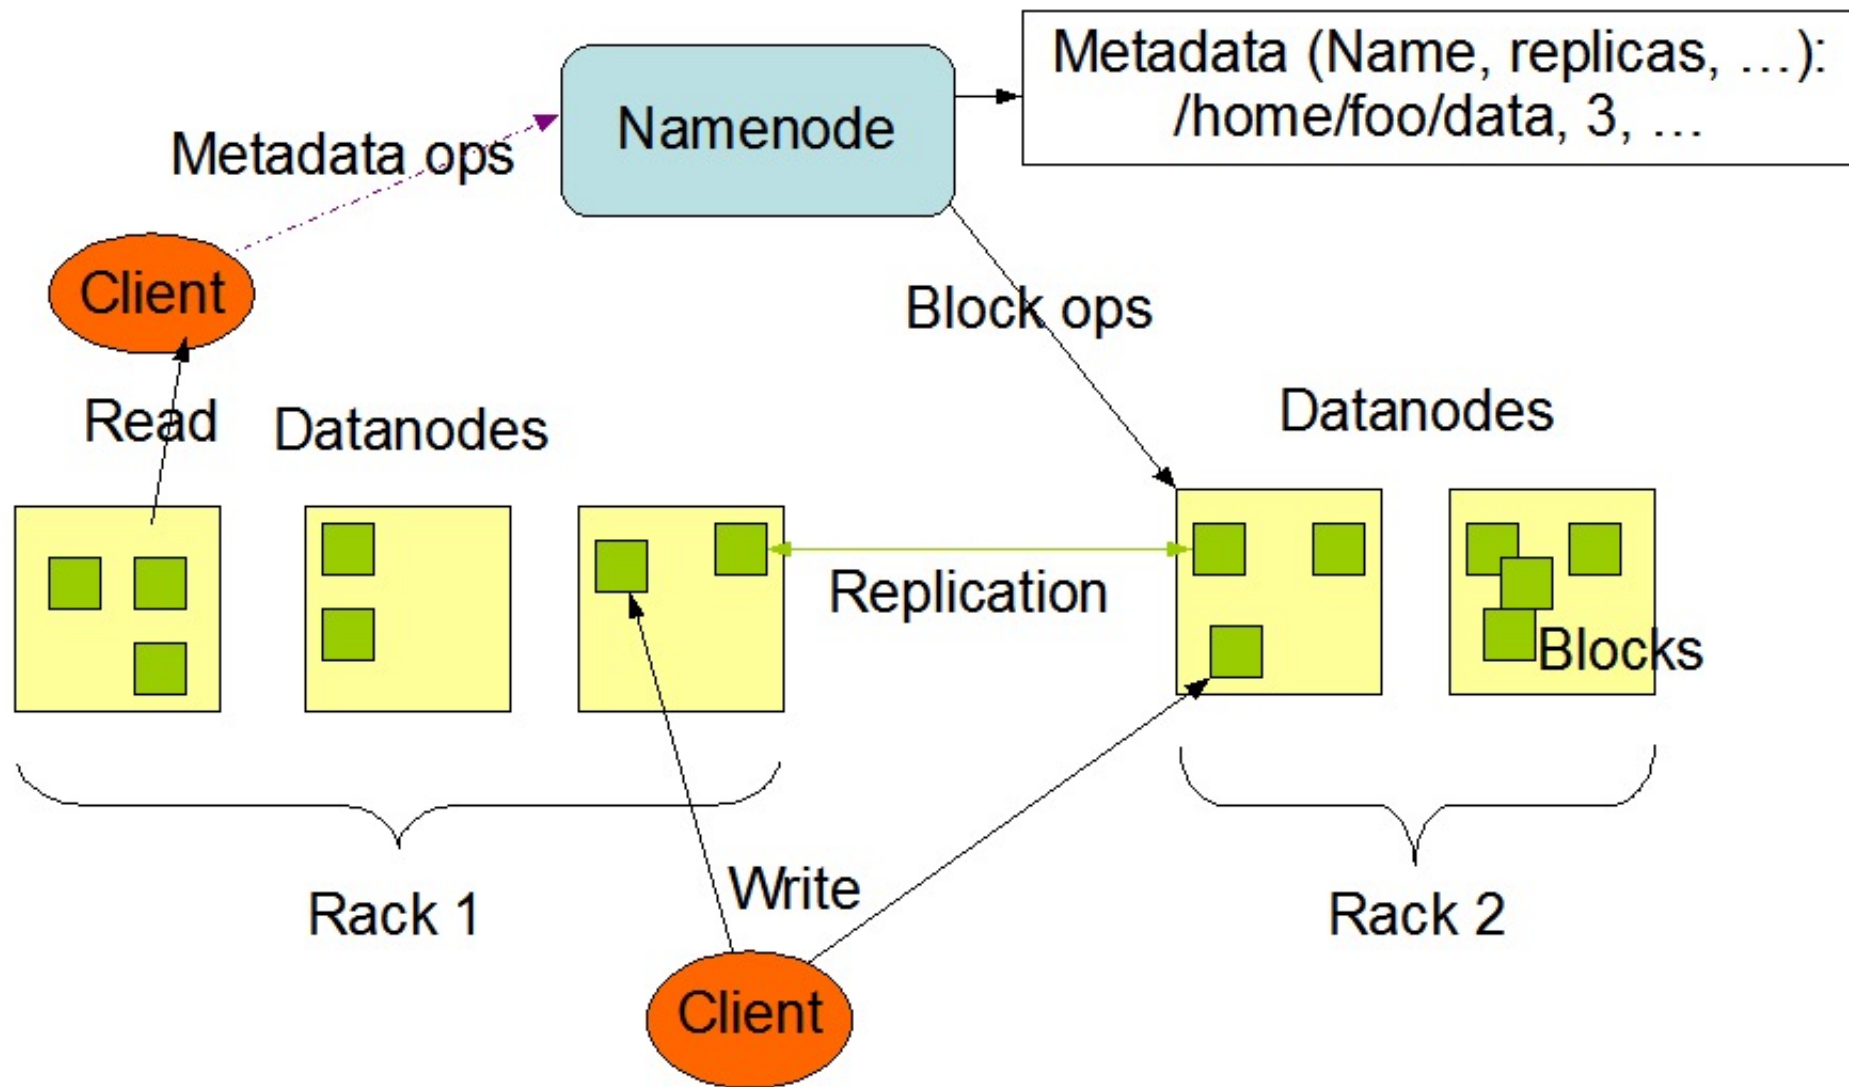


Architecture

- Data can be replicated
 - Replication factor (default 3)
 - More costly (space) than RAID or erasure codes
- Not all operations are supported
 - No random write, only append and truncate
- Permission
 - User and group for coarse grained control
 - Set access to owner or group of users
 - Access Control Lists (ACLs) for finer control
 - Set specific permission to specific user
 - Disabled by default

) → write once
Read Many

HDFS Architecture



Reading & writing

- All clients read and write requests go through NameNodes
 - Validation and metadata
- Actual data go directly to clients
 - Faster access
 - Blocks can be requested in parallel
- Reading workflow
 - Client send request for file “/test.txt”
 - NameNode checks access rights and returns list of blocks+DataNodes

Reading & writing

- Writing workflow
 - Client contact NameNode
 - If writing allowed, check if file exists
 - If yes, error
 - NameNode returns a list of DataNode
 - Client sends data to DataNodes in round-robin
 - After writing all blocks, the client notifies the NameNode
- Replication is handled by DataNodes while receiving data

HDFS Blocks

- Large blocks
 - Not suited for storing small files
 - Limit overhead of metadata
- Transfer cost of a block
 - Disk access + latency + throughput
 - Large blocks minimize impact of disk access and latency
- Replication for
 - Fault tolerance
 - Tries to put replica on different machine and different racks
 - Faster access
 - Load blocks from node closer to client

Network

Fault tolerance

- NameNode is a single point of failure
 - If down, cannot access data anymore
 - If destroyed (metadata lost), data are lost
- Possibility to use a Secondary NameNode
 - Maintains snapshot of metadata + edit log
 - Periodically apply edit log to metadata and store new state
- In case of crash
 - Restart NameNode
 - Get snapshot + log of Secondary NameNode and rebuilt recent state
- But
 - Rebuilding might be long
 - Might still lose some metadata

Fault tolerance

- *High Availability*
 - Feature introduced in Hadoop 2
- Support 2 NameNode
 - 1 active and 1 passive in standby
 - If active fails, standby takes its place
- How to ensure consistency?
 - Relies on JournalNode
 - Active write edit log to JournalNode
 - Passive regularly get edit log from JournalNode
- Protection against failures of JournalNode
 - Usually use 3 of them but N possible
 - Agreement based on a Quorum algorithm
 - Can tolerate $(N-1)/2$ failures